

## Logistic regression models for binary data

FW8051 Statistics for Ecologists

Department of Fisheries, Wildlife and Conservation Biology



### Learning objectives

- Be able to formulate, fit, and interpret logistic regression models appropriate for binary data using R and JAGS
- Be able to compare models and evaluate model fit
- Be able to visualize models using effect plots
- Be able to describe statistical models and their assumptions using equations and text and match parameters in these equations to estimates in computer output

## Logistic regression

Model for binary (0/1) data or binomial data (number of 1's out of  $n$  trials).

$$Y_i|X_i \sim \text{Binomial}(n_i, p_i)$$

$$\text{logit}(p_i) = \log\left(\frac{p_i}{1-p_i}\right) = \beta_0 + \beta_1 X_{1,i} + \dots + \beta_p X_{p,i}$$

- Random component = Bernoulli or binomial distribution
- Systematic component:  $\text{logit}(p_i)$  or  $\log(\text{odds})$  = linear combination of predictors

Remember, for binary data,  $E[Y_i|X_i] = p_i$ ,  $\text{Var}[Y_i|X_i] = p_i(1 - p_i)$

$$\Rightarrow p_i = \frac{\exp(\beta_0 + \beta_1 X_{1,i} + \dots + \beta_p X_{p,i})}{1 + \exp(\beta_0 + \beta_1 X_{1,i} + \dots + \beta_p X_{p,i})} \quad (\text{can use } \text{plogis} \text{ function in R})$$

## Logistic regression

$$Y_i|X_i \sim \text{Binomial}(n_i, p_i)$$

$$\text{logit}(p_i) = \log\left(\frac{p_i}{1-p_i}\right) = \beta_0 + \beta_1 X_{1,i} + \dots + \beta_p X_{p,i}$$

$\frac{p_i}{1-p_i}$  is referred to as the **odds**.

The link function,  $\log\left(\frac{p_i}{1-p_i}\right)$ , is referred to as **logit**.

Thus, we can describe our model in the following ways:

- We are modeling  $\log\left(\frac{p_i}{1-p_i}\right)$  as a linear function of  $X_1, \dots, X_p$ .
- We are modeling the logit of  $p$  as a linear function of  $X_1, \dots, X_p$ .
- We are modeling the log odds of  $p$  as a linear function of  $X_1, \dots, X_p$ .

$$\text{Odds} = \frac{p}{1-p}$$

If the probability of winning a bet is  $2/3$ , what are the odds of winning?

$$\text{odds} = \frac{p}{1-p} = (2/3 \div 1/3) = 2 \text{ (or "2 to 1")}$$

Table 6.1. Various probabilities, odds and log odds. The table shows how log odds are calculated from probabilities.

$p_i$	0.001	0.1	0.3	0.4	0.5	0.6	0.7	0.9	0.999
$1-p_i$	0.999	0.9	0.7	0.6	0.5	0.4	0.3	0.1	0.001
$O_i$	0.001	0.11	0.43	0.67	1	1.5	2.33	9	999
$\text{Ln}(O_i)$	-6.91	-2.20	-0.85	-0.41	0	0.41	0.85	2.20	6.91

From Zuur et al. 2007. Analyzing Ecological data

Odds can vary between 0 and  $\infty$ , so  $\log(\text{odds})$  can live on  $-\infty$  to  $\infty$ .

## Odds Ratios: $\exp(\beta)$

Consider a regression coefficient for a categorical variable:

$$\text{logit}(p_i) = \log\left(\frac{p_i}{1-p_i}\right) = \beta_0 + \beta_1 I(\text{group} = B)_i$$

$I(\text{group} = B)_i = 1$  if observation  $i$  is from Group B and 0 if Group A

- Odds for group B =  $\frac{p_B}{1-p_B} = \exp(\beta_0 + \beta_1)$
- Odds for group A =  $\frac{p_A}{1-p_A} = \exp(\beta_0)$

Consider the ratio of these odds:

$$\frac{\frac{p_B}{1-p_B}}{\frac{p_A}{1-p_A}} = \frac{\exp(\beta_0 + \beta_1)}{\exp(\beta_0)} = \frac{e^{\beta_0} e^{\beta_1}}{e^{\beta_0}} = \exp(\beta_1)$$

So,  $\exp(\beta_1)$  gives an **odds ratio** (or ratio of odds) for Group B relative to group A.

## Odds Ratios: $\exp(\beta)$

Consider a continuous predictor,  $X$ :

$$\text{logit}(p_i) = \log\left(\frac{p_i}{1-p_i}\right) = \beta_0 + \beta_1 X_i$$

$\beta_1$  gives the change in log odds per unit change in  $X$ .

- Odds when  $X_i = a$  is given by  $\frac{p_i}{1-p_i} = \exp(\beta_0 + \beta_1 a)$
- Odds when  $X_i = a + 1$  is given by  $\frac{p_i}{1-p_i} = \exp(\beta_0 + \beta_1(a + 1))$

Consider the ratio of these odds:

$$\frac{\frac{p_i}{1-p_i}}{\frac{p_i}{1-p_i}} = \frac{\exp(\beta_0 + \beta_1(a+1))}{\exp(\beta_0 + \beta_1 a)} = \frac{e^{\beta_0} e^{\beta_1 a} e^{\beta_1}}{e^{\beta_0} e^{\beta_1 a}} = \exp(\beta_1)$$

So,  $\exp(\beta_1)$ , gives the odds ratio for two observation that differ by 1 unit of  $X$ .

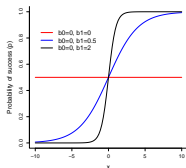
## Multiple predictors

For multiple predictor models,

$\exp(\beta_i)$  gives the odds ratio for observations where  $X_i$  differs by 1 unit, while holding everything else constant!

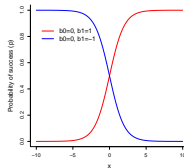
The odds is expected to increase by a factor of  $\exp(\beta_i)$  when  $X_i$  increases by 1 unit, and everything else is held constant!

$$\text{logit}(p_i) = \log\left(\frac{p_i}{1-p_i}\right) = \beta_0 + \beta_1 X_i$$



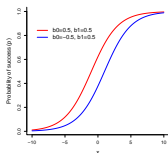
The slope coefficient  $\beta_1$  controls how quickly we transition from 0 to 1.

$$\text{logit}(p_i) = \log\left(\frac{p_i}{1-p_i}\right) = \beta_0 + \beta_1 X_i$$



The sign of  $\beta_1$  determines if  $p$  increases or decreases as we increase  $X$ .

$$\text{logit}(p_i) = \log\left(\frac{p_i}{1-p_i}\right) = \beta_0 + \beta_1 X_i$$



$\beta_0$ :

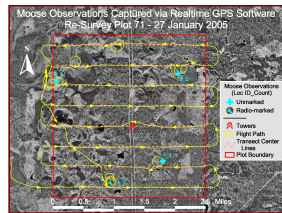
- Controls the height of the curve when  $X = 0$ .
- Gives the log odds of detection when all predictor variables = 0
- $E[Y_i|X_i = 0] = \frac{\exp(\beta_0)}{1 + \exp(\beta_0)}$  (equals 1/2 if  $\beta_0 = 0$ ).

## Sightability Surveys: Minnesota Moose

124 'trials', 2005-2007

$n_0 = 65$  missed groups

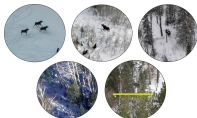
$n_1 = 59$  observed groups



- Binary observations,  $Y_i = 0$  (missed) or 1 (seen).
- Covariates thought to influence detection.

## Covariates

- Visual obstruction
- Survey year (may be due to different observers)



$$Y_i|X_i \sim \text{Bernoulli}(p_i)$$

$$\text{logit}(p_i) = \log\left(\frac{p_i}{1-p_i}\right) = \beta_0 + \beta_1 \text{voc}_i$$

Assumptions:

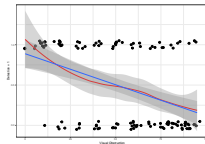
- observations are independent
- log odds is a linear function of voc
- mean and variance depend on voc

$$E[Y_i|X_i] = p_i; \text{Var}[Y_i|X_i] = p_i(1-p_i) \text{ with:}$$

$$p_i = \frac{\exp(\beta_0 + \beta_1 \text{voc}_i)}{1 + \exp(\beta_0 + \beta_1 \text{voc}_i)}$$

## Visual Obstruction

```
ggplot(exp.m, aes(voc, observed)) + theme_bw() +
  geom_point(position = position_jitter(w = 2, h = 0.05), size = 3) +
  geom_smooth(colour = "red") + geom_smooth(method = "lm") +
  xlab("Visual Obstruction") +
  ylab("Detection = 1")
```



- lm would eventually predict  $p_i \geq 1$  and  $p_i \leq 0$
- lm assumes constant variance rather than  $\text{var}(p_i) = p_i(1-p_i)$

```
mod1 <- glm(observed ~ voc, data = exp.m, family = binomial())
summary(mod1)
```

```
Call:
glm(formula = observed ~ voc, family = binomial(), data = exp.m)
```

```
Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.8056  -0.9071  -0.6218   0.9745   1.8647
```

```
Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  1.759933   0.460137   3.825 0.000131 ***
voc          -0.034792   0.007753  -4.487 7.21e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)
```

```
Null deviance: 171.61  on 123  degrees of freedom
Residual deviance: 147.38  on 122  degrees of freedom
AIC: 151.38
```

```
Number of Fisher Scoring iterations: 4
```

```
mod1$coef
```

```
(Intercept)      voc  
1.75993309 -0.03479153
```

Regression coefficient for voc (visual obstruction) = -0.039.

- The log odds of being detected decreases by 0.039 per unit increase in visual obstruction
- The odds of being detected decreases by a factor of  $\exp(0.039) = 0.96$  per unit increase in visual obstruction

Intercept = 2.12 =  $\log(\text{odds})$  of detection when VOC = 0.

```
# p(Y=1|voc=0) = exp(coef(mod1)[1]) / (1 + exp(coef(mod1)[1]))  
plogis(coef(mod1)[1])
```

```
(Intercept)  
0.8532013
```

We see roughly 90% of moose if there is no visual obstruction.

```
exp.m$year<-as.factor(exp.m$year)  
mod2<-glm(observed~voc+year, data=exp.m, family=binomial())  
summary(mod2)
```

```
Call:  
glm(formula = observed ~ voc + year, family = binomial(), data = exp.m)
```

```
Deviance Residuals:  
    Min       1Q   Median       3Q      Max  
-1.9351 -0.8411 -0.4561  0.9493  1.8680
```

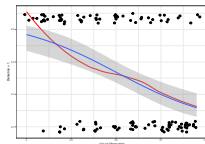
```
Coefficients:  
              Estimate Std. Error z value Pr(>|z|)  
(Intercept)  2.453203    0.622248   3.942 5.06e-05 ***  
voc          -0.037391    0.008199  -4.560 5.11e-06 ***  
year2006     -0.453862    0.516567  -0.879  0.3796  
year2007     -1.111884    0.508269  -2.188  0.0287 *  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 171.61  on 123  degrees of freedom  
Residual deviance: 142.23  on 120  degrees of freedom  
AIC: 150.23
```

Number of Fisher Scoring iterations: 4

```
ggplot(exp.m, aes(voc, observed)) + theme_bw() +  
  geom_point(position = position_jitter(w = 2, h = 0.05), size = 3) +  
  xlab("Visual Obstruction") + geom_smooth(se = F, colour = "red") +  
  stat_smooth(method = "glm", method.args = list(family = "binomial")) +  
  ylab("Detection = 1")
```



```
coef(mod2)
```

```
(Intercept)      voc      year2006      year2007  
2.45320264 -0.03739118 -0.45386154 -1.11188432
```

Year 2005:  $\log(p_i / (1 - p_i)) = 2.45 - 0.037VOC$

Year 2006:  $\log(p_i / (1 - p_i)) = 2.45 - 0.037VOC - 0.45$

So, -0.45 gives the difference in log odds between years 2005 and 2004 (if we hold VOC constant).

$\exp(-0.45) = 0.63$  = odds ratio (year 2006 to year 2005)

odds ratio =  $\frac{p_{2006} / (1 - p_{2006})}{p_{2005} / (1 - p_{2005})} = 0.63$

## Supporting Theory

The estimates of  $\beta$  are maximum likelihood estimates, found by maximizing:

$$L(\beta; y, x) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}, \text{ with}$$

$$p_i = \frac{e^{\beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik}}}{1 + e^{\beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik}}}$$

Remember, for large samples,  $\hat{\beta} \sim N(\beta, \Sigma)$ .

We can use this theory to conduct tests (z-statistics and p-values in output by the `summary` function) and to get confidence intervals.

- $\text{logit}(p) = \beta_0 + \beta_1 x_1 + \dots + \beta_k x_k$  is more "Normal" than  $p$
- Generate confidence intervals for  $\text{logit}(p)$ , then back-transform to get confidence intervals for  $p$
- Ensures the confidence intervals will live on the (0,1) scale
- Intervals will not be symmetric

## Confint

```
(ci.prof<-confint(mod2))
```

Waiting for profiling to be done...

	2.5 %	97.5 %
(Intercept)	1.30341777	3.7586692
voc	-0.05448153	-0.0221268
year2006	-1.48479529	0.5516852
year2007	-2.14380706	-0.1382692

These are profile-likelihood based confidence intervals based on "inverting" the likelihood ratio test (see Maximum Likelihood notes).

```
exp(ci.prof[,3,])
```

	2.5 %	97.5 %
	0.2265487	1.7361764

Profile-likelihood based intervals should have better statistical properties with small data sets (better **coverage** rates).

If confidence limits for  $\beta$  include 0 or confidence limits for  $\exp(\beta)$  include 1, then we do not have enough evidence to say that years differ in their detection probabilities.

```
mod2$coef
```

(Intercept)	voc	year2006	year2007
2.45320264	-0.03739118	-0.45386154	-1.11188432

```
sqrt(diag(vcov(mod2)))
```

(Intercept)	voc	year2006	year2007
0.622247867	0.008199483	0.516567443	0.508269279

```
exp(rep(0.006879, 2)+c(-1.96, 1.96)*0.53664) # exp(beta +/-1.96SE)
```

```
[1] 0.3517145 2.8826021
```

95% Confidence interval for odds ratio = (0.35, 2.88) includes 1 (not statistically significant)

## Goodness-of-fit

Can adapt our general approach for testing goodness-of-fit using Pearson residuals ( $r_i$ )

$$r_i = \frac{Y_i - E[Y_i|X_i]}{\sqrt{\text{Var}[Y_i|X_i]}}$$

- $E[Y_i|X_i] = p_i = \frac{\exp(\beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik})}{1 + \exp(\beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik})}$
- $\text{Var}[Y_i|X_i] = p_i(1 - p_i)$

See textbook for an implementation of this test...

## Hosmer-Lemeshow test (similar test)

Group Observations by deciles of their predicted values to form groups, then calculate the expected and observed number of successes and failures for each group:

Model Results				
	$G_1 = [0, \hat{\pi}_1]$	$G_2 = (\hat{\pi}_1, \hat{\pi}_2]$	...	$G_{g-1} = (\hat{\pi}_{g-1}, 1]$
Successes	$\sum_{i \in G_1} \hat{\pi}_i$	$\sum_{i \in G_2} \hat{\pi}_i$	...	$\sum_{i \in G_{g-1}} \hat{\pi}_i$
Failures	$n_1 - \sum_{i \in G_1} \hat{\pi}_i$	$n_2 - \sum_{i \in G_2} \hat{\pi}_i$	...	$n_{g-1} - \sum_{i \in G_{g-1}} \hat{\pi}_i$

Observed Results				
	$G_1 = [0, \hat{\pi}_1]$	$G_2 = (\hat{\pi}_1, \hat{\pi}_2]$	...	$G_{g-1} = (\hat{\pi}_{g-1}, 1]$
Successes	$\sum_{i \in G_1} Y_i$	$\sum_{i \in G_2} Y_i$	...	$\sum_{i \in G_{g-1}} Y_i$
Failures	$n_1 - \sum_{i \in G_1} Y_i$	$n_2 - \sum_{i \in G_2} Y_i$	...	$n_{g-1} - \sum_{i \in G_{g-1}} Y_i$

$$\chi^2 = \sum_{i=1}^{n_g} \frac{(O_i - E_i)^2}{E_i} \sim \chi_{g-2}^2$$

where  $g$  = number of groups.

See: **Goodness of fit with binary data** here: <http://www.unc.edu/courses/2010fall/ecol/563/001/docs/lectures/lecture21.htm>

## Hosmer-Lemeshow test

## Hosmer-Lemeshow Test

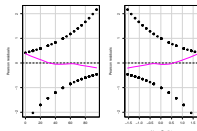
## Residual plots

```
library(ResourceSelection)
hoslem.test(exp.m$observed, fitted(mod1), g=8)
```

Hosmer and Lemeshow goodness of fit (GOF) test

```
data: exp.m$observed, fitted(mod1)
X-squared = 3.2505, df = 6, p-value = 0.7768
```

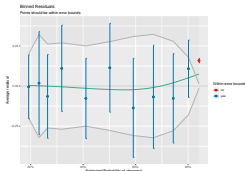
```
car::residualPlots(mod1)
```



```
Test stat Pr(>|Test stat|)
voc      0.6873      0.4071
```

## Binned residual plot

```
binplot<-performance::binned_residuals(mod1)
plot(binplot)
```



## Likelihood ratio tests

We can again use difference in deviances (equivalent to likelihood ratio tests) to compare full and reduced models.

```
drop1(mod2, test="Chisq")
```

Single term deletions

```
Model:
observed ~ voc + year
Df Deviance   AIC      LRT Pr(>Chi)
<none>          142.23 150.23
voc      1   168.20 174.20 25.9720 3.464e-07 ***
year     2   147.38 151.38  5.1558 0.07593 .
--- ---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

voc is an important predictor, the importance of year is less clear.

## ANOVA function (car package)

Or use Anova in car package

```
library(car)
```

```
Anova(mod2)
```

Analysis of Deviance Table (Type II tests)

```
Response: observed
      LR Chisq Df Pr(>Chisq)
voc    25.9720  1 3.464e-07 ***
year    5.1558  2 0.07593 .
--- ---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## AIC

We can compare nested or non-nested models using the AIC function

```
AIC(mod1, mod2)
```

	df	AIC
mod1	2	151.3824
mod2	4	150.2266



We can also summarize models by getting predicted values:  
 $P(\text{detect animal}|\text{voc})$ :

- $\text{logit}(p_i) = \beta_0 + \beta_1 x_1 + \dots + \beta_k x_k$
- $P(Y_i = 1|X = x) = p_i = \frac{e^{\beta_0 + \beta_1 x_1 + \dots + \beta_k x_k}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_k x_k}}$  (inverse logit)

We can use `predict(model, newdata=, type="link", se=TRUE)` to get predictions on logit scale.

Then use `plogis(p.hat$fit +/- 1.95*p.hat$se.fit)` to transform the limits back to the probability scale.

Model 2:  $\text{observed} \sim \text{voc} + \text{year}$  is additive on the logit scale

- Differences in  $\text{logit}(p)$  among years will not depend on  $\text{voc}$
- Differences in  $p$ , will however, depend on  $\text{voc}$ !

See: Section 16.6.3 in the book

- Can always create your own "effect" plots by calculating predicted values for different combinations of your predicted values
- Can use the `effects` package or `ggeffects` to do something similar

## Effect plots on probability scale

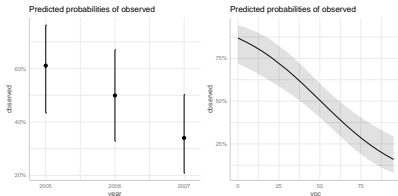
## Effect plots

Use `effects` or `ggeffects` package:

- Fixes all continuous covariates (other than the one of interest) at their mean values
- Categorical predictors: averages predictions on link scale, weighted by proportion of data in each category, then back transforms to probability scale
- These are referred to as marginal predictions by `ggeffects`

Use `type="response"` to plot on response scale

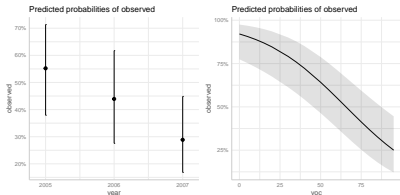
```
library(ggeffects); library(patchwork)
p1 <- plot(ggeffect(mod2, "year"))
p2 <- plot(ggeffect(mod2, "voc"))
p1 + p2
```



## Adjusted plots

Instead of averaging predictions across years, we could set year to a specific value. This leads to adjusted plots.

```
library(ggeffects); library(patchwork)
p1 <- plot(ggpredict(mod2, "year"))
p2 <- plot(ggpredict(mod2, "voc"))
p1 + p2
```



## JAGS

Will use a similar structure as we used for count models:

- A linear predictor,  $\eta = \beta_0 + \beta_1 x_1$  ( $x_1 = \text{voc}$ )
- $p_i = g^{-1}(\eta) = \frac{e^{\eta_i}}{1 + e^{\eta_i}}$
- $Y[i] \sim \text{dbin}(p[i], 1)$
- Require priors for  $\beta_0$  and  $\beta_1$ , e.g.,  $N(0, 0.01)$

Gelman's recommendations (see [arxiv.org/pdf/0901.4011.pdf](https://arxiv.org/pdf/0901.4011.pdf)):

- scale continuous predictors so they have mean 0 and sd = 0.5
- using a non-informative Cauchy prior  $\text{dt}(0, \text{pow}(2.5, -2), 1)$

In class exercise: adapt the JAGS code for fitting mod1 (voc only) to allow fitting of mod2 (voc + year).